

Development of an AI-Driven Criminal Face Detection and Identification System Using Hybrid Deep Learning Pipelines

Abstract

Dynamic face recognition in unconstrained surveillance environments remains one of the most critical challenges in computer vision and public safety engineering. Traditional biometric monitoring systems suffer from high false-positive rates, computational bottlenecks, and performance degradation under severe pose variations, fluctuating illumination, and partial facial occlusions. This paper outlines the academic proposal and technical blueprint for an automated criminal identification system designed using the hybrid Python framework *deepface*. By wrapping state-of-the-art deep convolutional neural networks—specifically the Additive Angular Margin Loss (ArcFace) model—and pairing them with high-performance detector backends such as RetinaFace, the proposed system establishes a robust pipeline encompassing real-time face detection, landmark-based alignment, high-dimensional embedding extraction, and rapid vector similarity matching. To address the scalability challenges associated with searching large criminal databases, the backend is integrated with a PostgreSQL vector database utilizing the *pgvector* extension to execute sub-second approximate nearest neighbor queries on millions of entries. Furthermore, the architecture is supplemented with passive liveness detection modules to prevent presentation attacks, alongside strict data protection and bias mitigation workflows. The resulting framework provides a highly accurate, computationally efficient, and ethically aligned decision-support tool for modern law enforcement and municipal security agencies.

1. Background and Contextual Framework

The rapid urbanization and increasing complexity of modern security landscapes have placed immense strain on conventional policing models, which rely heavily on manual monitoring, retrospective video analysis, and eyewitness testimonies.¹ Human observation under high-stress conditions is inherently prone to cognitive fatigue, perceptual errors, and sub-optimal response latency, leading to critical delays in locating fugitives or preventing active security breaches.² To alleviate these limitations, automated facial recognition technology (FRT) has emerged as an indispensable tool for public safety, enabling modern security infrastructure to transition from passive, post-event evidence collection to active, real-time situational awareness and threat detection.²

Early biometric identification methods operated on handcrafted spatial representations, such as Principal Component Analysis (PCA) for Eigenfaces or Local Binary Patterns Histograms (LBPH).¹ While computationally efficient, these classical algorithms require highly controlled

1	2026	<i>Crim eVisi on: An AI-Driven Face Reco gnition Fram ewor k for Crimi nal Identi ficati on</i>	He et al. ³	Propr ietary Susp ect DB; Live Surve illanc e Feed s	Frag ment ed data base s, slow manu al moni torin g, and com munit y publi c repor ting integ ratio n. ³	Deep Face with ArcFa ce; Retin aFac e & MTC NN back ends; FastA PI; React with Vite. ² 5	Syste m Laten cy, Preci sion, Geol ocati on Map ping Accu racy. ³	Achie ved robus t matc hing using fallin g back ends (Retin aFac e, MTC NN, Open CV) with a two-stage matc h valida tion. ² 5	Adap t the platfo rm to high-cong ested multi -face envir onme nts and distri bute d edge devic es. ⁵
2	2026	<i>Boos ted Light Face: A Hybri d DNN and GBM Mod el for Boos ted</i>	Seren gil & Özpi nar ¹⁵	Label ed Faces in the Wild (LFW); 6,00 0 imag e pairs. 24	Gene raliza tion error s and opti mizat ion limits of singl e deep conv	Ense mble of 5 DNN dista nce metri cs traine d via a Light GBM classi	Accu racy, Preci sion, Recal l, F1-Score, AUC. 24	Achie ved 99.1% LFW accur acy, outp erfor ming huma n baseli ne (97.5	Integ rate low-powe r edge opti mizat ion for the ense mble pipeli

		<i>Facial Recognition</i>			olutional back bones. ²²	fier. ²²		%) and Face Net512d (98.4%). ²⁴	ne. ²³
3	2025	<i>Deep Learning Approaches for Criminal Face Detection and Recognition: A Comparative Study</i>	Venkataratnam et al. ⁹	WIDER FACE ; Proprietary Suspect Database	Face detection in crowded areas , occlusion, and extreme pose variations. ⁹	Face Net (recognition); Faster R-CNN, YOLOv4, SSD (detection). ¹⁷	Bounding-box Recall, Precision, Inference FPS, Recognition Accuracy. ⁹	Faster R-CNN achieved superior occlusion accuracy but was slow. YOLOv4 was optimal for real-time video. ⁹	Combine YOLO models for bounding-box prediction with deep metric backbone s for classification. ⁹
4	2022	<i>Realtime Criminal Face Recognition</i>	Sanjay & Priya ²⁹	Saveetha Web Ontology Lab dataset	Sensitivity to lighting changes and	ArcFace over DeepFace; LBPH and classi	Verification Accuracy, Standard Devia	ArcFace achieved 92.50 % accuracy,	Evaluate performance variations across

		<i>System Based on Arcface over Deep face</i>			edge boundaries in unconstrained public areas. ²⁹	c CNN baselines. ²⁹	tion, Statistical Significance (p-value). ²⁹	outperforming Deep Face (91.90%) and LBPH (90.80%) with a significance of $p = 0.0118$. ²⁹	s multi-spectral or thermal imaging models. ⁵
5	2026	<i>Criminal Identification Using ResNet-Based Deep Pipelines</i>	Murillo-Moreno et al. ⁴	National Identity Database (NIN); Surveillance video	Verification failure of LBPH in unconstrained environments; first-time offender processing. ⁴	Haar Cascade (detection) + ResNet (recognition) vs. Haar + LBPH. ⁴	Precision, Matching Time, National Registry cross-reference rate. ⁴	ResNet achieved higher accuracy and faster matching times than traditional LBPH pipelines. ⁴	Develop lightweight mobile biometric apps to support field operations. ³³

These studies demonstrate a clear technological shift: traditional, handcrafted classifiers are being replaced by parameterized deep metric learning architectures.⁴ Models such as FaceNet

employ triplet loss functions to map faces into compact spaces where Euclidean distance directly represents identity similarity.¹⁰ ArcFace further improves class separation by applying additive angular margin constraints directly in the hypersphere space, which helps prevent feature overlap under occluded conditions.²⁰ To resolve the conflict between high model parameterization and real-time processing requirements, current research is actively focusing on hybrid ensemble systems (such as Boosted LightFace)²² and optimized indexing solutions on specialized vector databases.¹⁴

3. Core Research Questions

To systematically address the technical bottlenecks in unconstrained biometric surveillance, this project is structured around the following four research questions:

1. Under unconstrained surveillance conditions (including variable lighting, partial masking, and head pose variations), which deep representation backend wrapped by deepface (such as ArcFace, FaceNet, or VGG-Face) maximizes identification accuracy while maintaining sub-second verification latency?⁹
2. To what extent does integrating advanced face detection backends (such as RetinaFace or MTCNN) and facial landmark alignment reduce localized spatial noise and improve downstream identification performance compared to classical Haar Cascade models?⁴
3. How can high-dimensional vector search indices (such as pgvector on PostgreSQL or Hierarchical Navigable Small World graphs via FAISS) be integrated into a Python REST API to achieve sub-second matching latency across a criminal database containing more than one million entries?¹⁴
4. How does demographic representation bias (e.g., gender, race, age) affect off-the-shelf representation weights, and what threshold-calibration strategies can mitigate false matches for sensitive populations in law enforcement environments?²

4. Aim and Experimental Objectives

The primary aim of this academic project is to design, develop, and quantitatively evaluate a scalable, automated criminal face detection and real-time identification platform using the hybrid Python framework deepface.² The system is intended to convert static surveillance footage and public image reports into actionable intelligence for law enforcement agencies, enabling fast, precise, and secure matching against centralized databases of wanted individuals.¹⁸

To achieve this primary aim, the project is structured around the following four specific research objectives:

- **To analyze and optimize face detection co-usability:** Investigate and compare the performance of multiple detector backends (including RetinaFace, MTCNN, and SSD) within the deepface pipeline to determine the optimal configuration for fast, localized face bounding boxes under extreme camera angles and motion blur.¹¹
- **To design and tune the matching pipeline:** Develop an optimized facial verification pipeline using the deepface framework, evaluating representation models (such as ArcFace, FaceNet, and VGG-Face) under different distance metrics to identify the most

accurate configuration for occluded suspects.¹⁰

- **To implement high-scale vector indexing:** Integrate a relational database management system using PostgreSQL and the pgvector extension to run Approximate Nearest Neighbor (ANN) indexing, ensuring sub-second matching times at a scale of up to one million synthetic profiles.¹⁴
- **To evaluate system fairness and security:** Implement passive liveness detection and subgroup performance auditing to mitigate demographic biases while protecting data security through irreversible, non-invertible biometric vector embeddings.²⁶

5. Significance of the Study

The development of a highly accurate, real-time criminal face detection system has profound operational implications for public safety, criminal investigations, and smart city infrastructure.² Manual video analysis during active investigations is incredibly slow; hours of raw CCTV footage must be manually parsed to locate a suspect's path.¹ By implementing an automated deep learning pipeline, law enforcement can execute rapid, bulk database matches, transforming static cameras into proactive nodes that instantly alert field officers to the presence of wanted offenders.²

Furthermore, this study contributes significantly to the optimization of computer vision models deployed on edge computing systems.²³ Resolving the latency overhead of state-of-the-art DCNNs through hybrid pipelining, landmark-based alignment, and indexed databases allows complex AI models to run on resource-constrained hardware, such as mobile police units, body cameras, or localized security hubs.⁴

Importantly, this project actively addresses the ethical, legal, and racial biases that have historically marred the deployment of facial recognition technology in policing.² By integrating demographic-fairness testing³⁸, optimizing decision thresholds across subgroups², and aligning the software with rigid data protection frameworks—such as the UK Data Protection Act 2018, landmark judicial rulings (e.g., *Ed Bridges v. South Wales Police*), and the EU AI Act—the study establishes a baseline for trustworthy, legally compliant public biometrics.⁴¹

6. Boundaries of the Investigation and Project Limits

To ensure academic and technical feasibility, the operational boundaries of this project are defined across specific parameters:

In-Scope System Components

- The deployment of Python 3.10 and the deepface library to manage core face processing tasks, including face localization, alignment, representation extraction, and pairwise comparison.¹²
- The benchmarking of pre-trained deep representation networks (specifically ArcFace, FaceNet, and VGG-Face) alongside multiple face detection backends (RetinaFace, MTCNN, and SSD).⁹
- The configuration of a PostgreSQL relational database with the pgvector extension to store and index high-dimensional biometric vectors.¹⁴

- The generation of 100,000 synthetic vector profiles to simulate high-scale database query latency under Hierarchical Navigable Small World (HNSW) indexing.²¹
- The integration of passive single-frame liveness detection utilizing `anti_spoofing=True` in the deepface core pipeline to block presentation attacks.¹²
- The development of an end-to-end local web dashboard using FastAPI and React with Vite, featuring GPS mapping using Leaflet JS to display simulated suspect alerts.⁵

Out-of-Scope System Components

- Training deep convolutional neural network backbones from scratch, which would require massive high-performance computing clusters and millions of images.¹¹ Pre-trained weights optimized on web-scale datasets are utilized instead via transfer learning.¹¹
- Full 3D facial reconstruction, multi-spectral infrared sensor integration, or proprietary hardware modifications to physical CCTV cameras.⁶
- Autonomous law enforcement dispatching or physical intervention mechanisms; the system operates strictly as a decision-support platform, maintaining a critical "human-in-the-loop" verification step before any official intervention occurs.²

The boundary is drawn at this level to prioritize software engineering integration, algorithmic optimization, and backend database scalability over the creation of proprietary deep learning weights or physical sensor networks.¹

7. Execution Framework and Technical Procedural Pipeline

The proposed criminal identification system is structured as an end-to-end multi-stage pipeline, processing video frames and matching them against a secure database.

Core Processing Pipeline Stages


The pipeline executed by the backend is divided into five logical stages, managed internally by the deepface architecture:

1. **Detect:** Raw frames from video streams are captured using OpenCV.⁵ Faces are localized inside the frames, and bounding boxes are extracted.¹ The system defaults to RetinaFace as it outperforms other detectors on low-resolution and partially occluded images.¹¹
2. **Align:** Faces are spatially aligned based on localized facial landmark coordinates (specifically eye positions) to eliminate head tilt, rotation, and skew.¹¹ Landmark alignment increases downstream recognition accuracy by up to 6%.¹⁶
3. **Normalize:** Aligned facial regions are resized to the precise dimensions expected by the deep learning backbone (e.g., 112×112 pixels for ArcFace or 224×224 pixels for FaceNet and VGG-Face).¹² Pixel intensity values are scaled using base normalization to reduce covariate shift and equalize different lighting distributions.⁴
4. **Represent:** The normalized facial image is fed into a deep convolutional neural

network.¹⁰ The model acts as a feature extractor, projecting the spatial image data into a compact, highly discriminative, multi-dimensional vector embedding.⁷

5. **Verify / Search:** The generated vector embedding is queried against the registered suspect embeddings in the database to identify matches.¹⁴

Technical System Architecture & Workflow Flowchart



Mathematical Metric Formulation

During the verification stage, the similarity of two facial embedding vectors u and v in a d -dimensional vector space is determined using standardized distance metrics.¹⁰

Cosine Distance

The default distance metric configured in deepface is Cosine Distance (D_C), which measures the angular difference between two non-zero vectors regardless of their magnitude:

$$D_C(u, v) = 1 - \frac{u \cdot v}{\|u\|_2 \|v\|_2} = 1 - \frac{\sum_{i=1}^d u_i v_i}{\sqrt{\sum_{i=1}^d u_i^2} \sqrt{\sum_{i=1}^d v_i^2}}$$

Euclidean Distance (L_2 Normalized)

To achieve greater robustness to scaling variations, L_2 normalization is applied to the vectors

before calculating the Euclidean Distance (D_{L2}), ensuring that vectors are projected onto a hypersphere:

$$\hat{u} = \frac{u}{\|u\|_2}, \quad \hat{v} = \frac{v}{\|v\|_2}$$

$$D_{L2}(u, v) = \sqrt{\sum_{i=1}^d (\hat{u}_i - \hat{v}_i)^2}$$

If the calculated distance falls below a pre-tuned threshold (T), the system classifies the face pair as a match (i.e., identifying the suspect).¹⁰

Feature Comparison of Embedded DeepFace Models

The system allows dynamic switching of the underlying representation model. The table below summarizes the key features of the primary models evaluated for this study:

Model Backbone	Vector Dimensions (d)	Core Optimization Metric / Loss Function	Inherent Structural Advantages	Reference
VGG-Face	4096	Softmax / Cross-Entropy	Deep, highly parameterized representation; excellent for high-quality frontal portraits. ¹¹	9
FaceNet	128 or 512	Triplet Loss (minimizes intra-class variance)	Very compact embedding; highly efficient for memory storage and fast exact queries. ¹⁰	9
ArcFace	512	Additive Angular Margin	Maximizes geodesic	20

		Loss	distance on a hypersphere; outstanding robustness under occlusions. ²⁹	
SFace	512	Sigmoid-Const rained Hypersphere	Exceptional robustness to degraded, low-resolution surveillance footage. ¹⁹	15

Relational Database Vector Schema

A PostgreSQL database enabled with the pgvector extension is used to execute Approximate Nearest Neighbor (ANN) searches via a Hierarchical Navigable Small World (HNSW) index.¹⁴ The database table schema is configured as:

```
CREATE EXTENSION IF NOT EXISTS vector;
```

```
CREATE TABLE suspect_profiles (
  id SERIAL PRIMARY KEY,
  suspect_name VARCHAR(100) NOT NULL,
  alias VARCHAR(100),
  demographics JSONB,
  face_embedding vector(512) NOT NULL
);
```

```
CREATE INDEX suspect_embedding_hnsw_idx
ON suspect_profiles USING hnsw (face_embedding vector_cosine_ops);
```

Using this setup, query vectors are compared against stored profiles in milliseconds:

```
SELECT suspect_name, alias, face_embedding <=> $1 AS distance
FROM suspect_profiles
ORDER BY distance LIMIT 1;
```

(Where <=> denotes the Cosine Distance operator in pgvector).

If the returned minimum distance is below the optimal pre-tuned validation threshold of **0.58**, the system flags a positive suspect identification.¹⁰ If the query is from a citizen reporting an

incident, the system captures their GPS coordinates, updates a Vite-powered React interactive Leaflet map, logs a secure, non-modifiable entry into an audit table, and propagates immediate visual and audio alerts to the administrative dashboard.⁵

8. Requirements Resources

To build, test, and run the proposed criminal face detection and identification platform, the following hardware and software resources are required:

Hardware Infrastructure

Hardware Component	Technical Specification	Operational Role inside Project
System Processor	Intel Core i7-12700K (12 Cores, up to 5.0 GHz)	Manages the main application thread, REST API requests, and coordinate database queries. ²⁵
Graphics Accelerator	Nvidia GeForce RTX 3070 (8GB GDDR6, 5888 CUDA Cores)	Accelerates deep convolutional model inference and handles multi-frame video decoding. ²⁶
System Memory	32 GB DDR4 RAM (3200 MHz, Dual Channel)	Prevents memory bottlenecks when loading multiple pre-trained DCNN weights in parallel. ⁴⁶
Local Storage	1 TB NVMe M.2 SSD (Read speeds up to 3500 MB/s)	Supports fast write operations for high-frame-rate video files and quick database lookups. ⁴⁶
Video Capture Device	USB HD Web Camera (1080p, 30 Frames Per Second)	Simulates continuous live video streams from public security cameras. ⁵

Software Stack and Packages


Software Component	Target Version	Operational Role inside Project
Operating System	Ubuntu 22.04 LTS (64-bit)	Provides a stable Linux environment for containerized development and deployment. ⁴⁶
Runtime Environment	Python 3.10.12	Core programming environment to build backend APIs and process image vectors. ¹²
Facial Analysis Framework	deepface (v0.0.100 or higher)	Handles face detection, landmark alignment, and embedding extraction. ¹²
Image Processing Library	OpenCV (v4.7.0)	Manages video streaming, frame extraction, and spatial scaling. ⁵
Machine Learning Backend	TensorFlow (v2.12.0)	Serves as the deep learning backend for running pre-trained model weights. ⁵
Database Server	PostgreSQL (v15.3)	Stores facial database entries and runs fast vector similarity queries. ¹⁴
Vector Extension	pgvector (v0.4.1)	Enables HNSW indexing for fast approximate nearest neighbor (ANN) searches. ¹⁴
Backend REST API	FastAPI (v0.110.0 or higher)	Exposes asynchronous endpoints for high-throughput image uploads, verification, and alerts. ²⁵
Web Frontend	React with Vite (v18.3 / v5.x)	Renders the fast, lightweight administrative dashboard, live camera

feeds, and real-time maps.¹

9. Research Plan, Milestones, and Contingencies

The project is scheduled for completion over a 16-week timeline, divided into logical engineering phases to ensure successful delivery.

Operational Project Gantt Chart



Detailed Phase Execution

- **Weeks 1–2: Literature Analysis & Workspace Configuration**
 - Review key literature regarding deep face matching and unconstrained biometric surveillance.⁴
 - Set up the Python virtual environment and install OpenCV, TensorFlow, PyTorch, and deepface.¹²
- **Weeks 3–4: Database Engineering & Schema Design**
 - Install PostgreSQL and configure the pgvector database extension.²¹
 - Generate a secure database schema to store suspect profiles and write scripts to load synthetic profiles for scale testing.²¹
- **Weeks 5–7: Pipeline Tuning & Optimization**
 - Benchmark different detector backends (RetinaFace, MTCNN, SSD) and representation models (ArcFace, FaceNet).¹⁶
 - Define precise mathematical distance thresholds (T) to optimize the trade-off between False Acceptance Rates and False Rejection Rates.¹⁰
- **Weeks 8–10: Web Application & API Development**
 - Develop FastAPI REST APIs to handle registration, image validation, and live camera frame searches.⁵
 - Build the Vite-powered React frontend dashboard, integrating a Leaflet map to track alerts in real time.²⁵
- **Weeks 11–12: Validation, Performance Evaluation, & Bias Audits**
 - Test matching accuracy under simulated real-world degradations (low lighting, motion blur, occlusions).⁵
 - Evaluate database query response times as records scale to 100,000 synthetic profiles.²¹
 - Conduct fairness audits to identify performance differences across age, gender, and racial subgroups.³⁸
- **Weeks 13–14: Security Hardening & Vulnerability Auditing**

- Implement passive liveness detection and active challenge-response protocols to protect against spoofing.¹²
- Integrate token-based authentication and secure audit trails to safeguard access to criminal databases.²⁵
- **Weeks 15–16: Technical Documentation & Final Review**
 - Analyze experimental results, compile performance graphs, and draft the final project submission paper.⁵²
 - Package the application in Docker containers and deliver the final codebase to the academic committee.⁴⁷

Risk and Contingency Framework

Identified Operational Risk	Potential System Impact	Algorithmic / Engineering Mitigation Strategy
Lighting Fluctuation and Occlusions	Dramatic increase in False Rejection Rates (FRR) under poor ambient lighting. ⁵	Configure RetinaFace as the primary detector and MTCNN as a fallback; apply adaptive histogram equalization (AHE) to normalize contrast. ¹⁹
High Search Latency at Scale	Query matching times exceed 1 second when database grows to one million suspect profiles. ¹⁴	Transition PostgreSQL from linear exact searches to Approximate Nearest Neighbor (ANN) indexing via pgvector using HNSW. ¹⁴
Demographic Bias in Pre-trained Weights	Disproportionately high rates of false-positive matches for specific ethnic subgroups. ²	Establish dynamic, group-specific distance thresholds (T_{group}) calibrated using fairness-aware validation datasets. ²
Presentation Attacks (Spoofing)	Unauthorized individuals bypass system checks using printed photographs or digital screens. ⁴⁰	Enable anti_spoofing=True in the deepface backend to activate real-time passive liveness checks. ¹²
Data Breaches of	Intercepted database	Enforce localized face

Biometric Records	entries compromise sensitive facial images of private citizens. ⁴⁵	processing on secure servers, store only encrypted 512-dimensional vector embeddings, and encrypt database connections. ²⁶
--------------------------	---	---

Conclusions

The academic minor project proposed in this paper outlines a secure, scalable, and ethically aligned criminal face detection and real-time identification system using the hybrid Python library deepface.² By integrating advanced deep learning models (such as ArcFace and RetinaFace) with a PostgreSQL vector database powered by the pgvector extension, the proposed architecture bridges the gap between high-dimensional representation accuracy and sub-second query speeds.¹⁴

The modular technical design ensures the system is highly versatile, supporting exact local search indexing for smaller databases as well as fast approximate nearest neighbor (ANN) searches for large-scale operations.¹⁴ The inclusion of liveness detection modules, demographic fairness checks, and irreversible biometric embedding storage addresses critical data protection requirements, aligning the platform with modern ethical and legal standards.⁴⁰ Ultimately, this system delivers an advanced, secure, and highly reliable decision-support framework that significantly improves public safety while strictly safeguarding individual privacy rights.¹

Works cited

1. CRIMDETECT : AN AI-POWERED CRIMINAL FACE DETECTION AND IDENTIFICATION SYSTEM USING MACHINE LEARNING - IJNRD.org, accessed May 21, 2026, <https://www.ijnrd.org/papers/IJNRD2604269.pdf>
2. Automated Criminal Identification System Using Advanced Facial Recognition Techniques, accessed May 21, 2026, <https://www.computer.org/csdl/proceedings-article/assic/2025/11158654/2a5e0qX0z04>
3. Development and Implementation of Face Recognition Technology for Criminal Identification and Public Safety - AIJMR, accessed May 21, 2026, <https://www.aijmr.com/research-paper.php?id=1275>
4. AI based criminal identification system using facial recognition - World Journal of Advanced Research and Reviews, accessed May 21, 2026, https://wjarr.com/sites/default/files/fulltext_pdf/WJARR-2026-0569.pdf
5. Criminal Detection Using Real-Time Visual Recognition and Alert System Integrating Deep Learning Models and Surveillance Feeds f - IJSDR, accessed May 21, 2026, <https://ijsdr.org/papers/IJSDR2504226.pdf>
6. Automated Criminal Identification System Using Advanced Facial Recognition Techniques, accessed May 21, 2026,

- <https://ieeexplore.ieee.org/document/11158654/>
7. Criminal Face Detection System - Research Square, accessed May 21, 2026, https://assets-eu.researchsquare.com/files/rs-2605399/v2_covered_81c95303-030b-4fff-af75-0828e707eaaa.pdf
 8. AI-DRIVEN CRIMINAL IDENTIFICATION SYSTEM USING DEEP LEARNING AND FACIAL RECOGNITION FOR ENHANCED LAW ENFORCEMENT - IJETCSE, accessed May 21, 2026, <https://ijetcse.com/wp-content/uploads/AI-DRIVEN-CRIMINAL-IDENTIFICATION-SYSTEM-USING-DEEP-LEARNING-AND-FACIAL-RECOGNITION-FOR-ENHANCED-LAW-ENFORCEMENT.pdf>
 9. Deep Learning Approaches for Criminal Face Detection and Recognition: A Comparative Study - International Journal of Computer Applications, accessed May 21, 2026, <https://www.ijcaonline.org/archives/volume186/number72/venkataratnam-2025-ijca-924570.pdf>
 10. ADVANCE CRIMINAL DETECTION SYSTEM USING DEEP LEARNING TECHNIQUE - TIJER, accessed May 21, 2026, <https://tijer.org/tijer/papers/TIJER2505113.pdf>
 11. Face Recognition with Facebook DeepFace in Keras - Sefik Ilkin Serengil, accessed May 21, 2026, <https://sefiks.com/2020/02/17/face-recognition-with-facebook-deepface-in-keras/>
 12. serengil/deepface: A Lightweight Face Recognition and ... - GitHub, accessed May 21, 2026, <https://github.com/serengil/deepface>
 13. A Benchmark of Facial Recognition Pipelines and Co-Usability Performances of Modules, accessed May 21, 2026, <https://dergipark.org.tr/en/pub/gazibtd/issue/84331/1399077>
 14. Introducing Brand New Face Recognition in DeepFace - Sefik Ilkin Serengil, accessed May 21, 2026, <https://sefiks.com/2026/01/01/introducing-brand-new-face-recognition-in-deepface/>
 15. Boosted LightFace: A Hybrid DNN and GBM Model for Boosted Facial Recognition - Gazi University Journal of Science - DergiPark, accessed May 21, 2026, <https://dergipark.org.tr/en/pub/gujs/article/1794891>
 16. deepface-fork - PyPI, accessed May 21, 2026, <https://pypi.org/project/deepface-fork/>
 17. Deep Learning Approaches for Criminal Face Detection and Recognition: A Comparative Study - International Journal of Computer Applications, accessed May 21, 2026, <https://www.ijcaonline.org/archives/volume186/number72/deep-learning-approaches-for-criminal-face-detection-and-recognition-a-comparative-study/>
 18. Criminal Face Recognition System Using Artificial Intelligence - IJSART, accessed May 21, 2026, https://ijsart.com/public/storage/paper/pdf_2/IJSARTV12I1104533.pdf
 19. Robust Face Recognition Under Challenging Conditions: A Comprehensive Review of Deep Learning Methods and Challenges - MDPI, accessed May 21,

- 2026, <https://www.mdpi.com/2076-3417/15/17/9390>
20. Multi scale self attention fusion for facial recognition: robustness improvement in occluded scenes - Research Square, accessed May 21, 2026, https://assets-eu.researchsquare.com/files/rs-9398126/v1_covered_96151562-dfa5-4f9a-b6b4-e8bdb84b6258.pdf
 21. Postgres As A Vector Database: Billion-Scale Vector Similarity Search With pgvector, accessed May 21, 2026, <https://sefiks.com/2024/07/05/postgres-as-a-vector-database-billion-scale-vector-similarity-search-with-pgvector/>
 22. Boosted LightFace: A Hybrid DNN and GBM Model for Boosted Facial Recognition - DergiPark, accessed May 21, 2026, <https://dergipark.org.tr/en/download/article-file/5291132>
 23. Real-Time Multi-Modal Embedded Vision Framework for Object Detection Facial Emotion Recognition and Biometric Identification on Low-Power Edge Platforms - arXiv, accessed May 21, 2026, <https://arxiv.org/html/2601.11970v1>
 24. Boosted LightFace: A Hybrid DNN and GBM Model for Boosted Facial Recognition, accessed May 21, 2026, https://www.researchgate.net/publication/400362373_Boosted_LightFace_A_Hybrid_DNN_and_GBM_Model_for_Boosted_Facial_Recognition
 25. DEVELOPMENT AND IMPLEMENTATION OF FACE RECOGNITION TECHNOLOGY FOR CRIMINAL IDENTIFICATION AND PUBLIC SAFETY - AIJMR, accessed May 21, 2026, <https://www.aijmr.com/papers/2026/2/1275.pdf>
 26. Mr. Sefik Ilkin Serengil | Author - SciProfiles, accessed May 21, 2026, <https://sciprofiles.com/profile/4261407>
 27. A Benchmark of Facial Recognition Pipelines and Co-Usability Performances of Modules Yüz Tanıma Üretim Hatlarının ve Modü - DergiPark, accessed May 21, 2026, <https://dergipark.org.tr/en/download/article-file/3573195>
 28. Demystifying the Face Recognition Process: Algorithms and Models - Embien Technologies, accessed May 21, 2026, <https://www.embien.com/blog/demystifying-the-face-recognition-process-algorithms-and-models>
 29. Realtime Criminal Face Recognition System based on Innovative Facial Classification to Improve Accuracy using Arcface over Deepface, accessed May 21, 2026, <https://www.pnrjournal.com/index.php/home/article/download/935/735/1114>
 30. Realtime Criminal Face Recognition System based on Innovative Facial Classification to Improve Accuracy using Arcface over Deepface | Journal of Pharmaceutical Negative Results, accessed May 21, 2026, <https://www.pnrjournal.com/index.php/home/article/view/935>
 31. Realtime Criminal Face Recognition System based on Innovative Facial Classification to Improve Accuracy using Arcface over Deepface | Semantic Scholar, accessed May 21, 2026, <https://www.semanticscholar.org/paper/Realtime-Criminal-Face-Recognition-System-based-on-Sanjay-Priya/f8741f07ef3923c52f11526b942a17bca79d2bbc>
 32. Improved Accuracy of Face Recognition System to Identify Criminal Based on

- Innovative Feature Extractor to Improve Accuracy Using Deepface over LBPH Algorithm - Journal of Pharmaceutical Negative Results, accessed May 21, 2026, <https://www.pnrjournal.com/index.php/home/article/download/934/734/1113>
33. Detecting of Criminals using Facial Recognition - International Journal of Engineering Research & Technology, accessed May 21, 2026, <https://www.ijert.org/research/detecting-of-criminals-using-facial-recognition-IJERTV13IS100122.pdf>
 34. Surveillance System for Real-Time High-Precision Recognition of Criminal Faces From Wild Videos - IEEE Xplore, accessed May 21, 2026, <http://ieeexplore.ieee.org/iel7/6287639/10005208/10143202.pdf>
 35. deepface/deepface/DeepFace.py at master · serengil/deepface - GitHub, accessed May 21, 2026, <https://github.com/serengil/deepface/blob/master/deepface/DeepFace.py>
 36. A Benchmark of Facial Recognition Pipelines and Co-Usability Performances of Modules, accessed May 21, 2026, <https://www.semanticscholar.org/paper/A-Benchmark-of-Facial-Recognition-Pipelines-and-of-Serengil-%C3%96zp%C4%B1nar/4e5d8c9a1db8c1f6243e14e851713370a99b6cec>
 37. Is there a way to use deepface.find() with a db being a S3 or a db table containing embeddings? - Stack Overflow, accessed May 21, 2026, <https://stackoverflow.com/questions/79482581/is-there-a-way-to-use-deepface-find-with-a-db-being-a-s3-or-a-db-table-contain>
 38. LightFace: A Hybrid Deep Face Recognition Framework | Request PDF - ResearchGate, accessed May 21, 2026, https://www.researchgate.net/publication/346308816_LightFace_A_Hybrid_Deep_Face_Recognition_Framework
 39. A Forensic Without the Science: Face Recognition in U.S. Criminal Investigations - Georgetown Law, accessed May 21, 2026, <https://www.law.georgetown.edu/privacy-technology-center/publications/a-forensic-without-the-science-face-recognition-in-u-s-criminal-investigations/>
 40. Trustworthy Face Recognition as a Service: A Multi-Layered Approach for Mitigating Spoofing and Ensuring System Integrity - MDPI, accessed May 21, 2026, <https://www.mdpi.com/1999-5903/17/10/450>
 41. Briefing note on the ethical issues arising from public-private collaboration in the use of live facial recognition technology (accessible) - GOV.UK, accessed May 21, 2026, <https://www.gov.uk/government/publications/public-private-use-of-live-facial-recognition-technology-ethical-issues/briefing-note-on-the-ethical-issues-arising-from-public-private-collaboration-in-the-use-of-live-facial-recognition-technology-accessible>
 42. The ethics of facial recognition technologies, surveillance, and accountability in an age of artificial intelligence: a comparative analysis of US, EU, and UK regulatory frameworks - PMC, accessed May 21, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC8320316/>
 43. Enhanced Face Recognition Models: Leveraging the Power of Machine Learning

- and Deep Learning for Accurate Identification - IEEE Xplore, accessed May 21, 2026, <https://ieeexplore.ieee.org/document/11434136/>
44. Face Recognition System for Criminal Identification in CCTV Footage Using Keras and OpenCV - IJETA, accessed May 21, 2026, <https://www.ijeta.org/download/file/fid/163617>
 45. Full article: Facial recognition on trial: data protection, discrimination, and the ethics of algorithmic governance in policing - Taylor & Francis, accessed May 21, 2026, <https://www.tandfonline.com/doi/full/10.1080/13600834.2026.2638629>
 46. Facial Recognition and Crime Detection System, accessed May 21, 2026, https://courses.ece.cornell.edu/ece5990/ECE5725_Fall2019_Projects/Dec_05_Demo/Crime_Detection/zl766_yz2583_120519/zl766_yz2583_120519/index.html
 47. Releases · serengil/deepface - GitHub, accessed May 21, 2026, <https://github.com/serengil/deepface/releases>
 48. serengil/deepface-react-ui: Deep Face Recognition UI With ReactJS - GitHub, accessed May 21, 2026, <https://github.com/serengil/deepface-react-ui>
 49. Impact of Image Resolution on Age Estimation with DeepFace and InsightFace - arXiv, accessed May 21, 2026, <https://arxiv.org/html/2511.14689v1>
 50. (PDF) A Benchmark of Facial Recognition Pipelines and Co-Usability Performances of Modules - ResearchGate, accessed May 21, 2026, https://www.researchgate.net/publication/380180104_A_Benchmark_of_Facial_Recognition_Pipelines_and_Co-Usability_Performances_of_Modules
 51. Enhanced Criminal Identification Using Integrated Face Recognition and Blockchain Technology - Fringe Global Scientific Press, accessed May 21, 2026, <https://fringe-global.com/ojs/index.php/jsihs/article/download/enhanced-criminal-identification-using-integrated-face-recognition/166/870>
 52. Minor Project-Assignment 1.pdf
 53. The ethics of facial recognition in police practice: beyond the usual concerns, accessed May 21, 2026, <https://www.emerald.com/jcrpp/article/11/2/166/1263785/The-ethics-of-facial-recognition-in-police>